

◆ Cell 1: Install Required Libraries

```
!pip install -q bitsandbytes accelerate datasets loralib peft transformers trl
```

```
!pip install datasets
```

📌 Explanation:

Installs all the required libraries to run and fine-tune a QLoRA model using TinyLlama:

- bitsandbytes for 4-bit quantization.
 - accelerate, transformers, trl for Hugging Face model management and training.
 - loralib, peft for applying Low-Rank Adaptation (LoRA).
 - datasets to load the training dataset from Hugging Face Hub.
-

◆ Cell 2: Import Libraries & Load Base Model

```
import torch
```

```
from datasets import Dataset
```

```
from transformers import AutoTokenizer, AutoModelForCausalLM,  
TrainingArguments, Trainer, BitsAndBytesConfig
```

```
from peft import LoraConfig, get_peft_model,  
prepare_model_for_kbit_training
```

📌 Explanation:

Imports necessary modules to:

- Load the model & tokenizer
 - Configure LoRA and quantization
 - Setup training arguments
-

◆ Cell 3: Configure 4-bit Quantization (QLoRA)

```
model_name = "TinyLlama/TinyLlama-1.1B-Chat-v1.0"
```

```
bnb_config = BitsAndBytesConfig(
    load_in_4bit=True,
    bnb_4bit_compute_dtype=torch.float16,
    bnb_4bit_use_double_quant=True,
    bnb_4bit_quant_type="nf4",
)
```

✦ **Explanation:**

Loads the TinyLlama model with **4-bit quantization** using BitsAndBytesConfig, which reduces memory and speeds up training (QLoRA technique).

◆ **Cell 4: Load Pretrained Model with Quantization Config**

```
model = AutoModelForCausalLM.from_pretrained(
    model_name,
    quantization_config=bnb_config,
    device_map="auto"
)
```

✦ **Explanation:**

Downloads the pretrained TinyLlama model and applies the quantization configuration. `device_map="auto"` places the model across GPUs/CPU as needed.

◆ **Cell 5: Load Tokenizer & Prepare Model for LoRA**

```
tokenizer = AutoTokenizer.from_pretrained(model_name, use_fast=True)
tokenizer.pad_token = tokenizer.eos_token
model.config.use_cache = False

model.gradient_checkpointing_enable()
```

```
model = prepare_model_for_kbit_training(model)
```

✦ **Explanation:**

- Loads tokenizer and adjusts padding token.
 - Enables gradient_checkpointing to reduce memory.
 - Prepares the model for training with LoRA + quantization.
-

◆ **Cell 6: Configure & Apply LoRA**

```
lora_config = LoraConfig(  
    r=8,  
    lora_alpha=32,  
    target_modules=["q_proj", "v_proj"],  
    lora_dropout=0.05,  
    bias="none",  
    task_type="CAUSAL_LM"  
)  
  
model = get_peft_model(model, lora_config)
```

✦ **Explanation:**

Defines LoRA config targeting the attention projections (q_proj, v_proj) and applies LoRA to the model using get_peft_model.

◆ **Cell 7: Load Dataset**

```
from datasets import load_dataset  
  
dataset = load_dataset("MakTek/Customer_support_faqs_dataset",  
    split="train")  
  
print("Dataset columns:", dataset.column_names)
```

Explanation:

Loads the MakTek/Customer_support_faqs_dataset from Hugging Face, which contains support-related questions and answers.

◆ **Cell 8: Format Instructions as Prompt-Response Pairs**

```
def format_instruction(example):  
    return f"### Instruction:\n{example['question']}\n\n###  
Response:\n{example['answer']}"  
  
dataset = dataset.map(lambda x: {"text": format_instruction(x)})
```

Explanation:

Reformats each row of the dataset into an instruction-tuned format (helpful for chat-based models):

```
### Instruction:  
  
<user_question>  
  
  
### Response:  
  
<model_answer>
```

◆ **Cell 9: Tokenize the Dataset**

```
def tokenize_function(example):  
    tokenized = tokenizer(example["text"], truncation=True,  
padding="max_length", max_length=512)  
    tokenized["labels"] = tokenized["input_ids"].copy()  
    return tokenized  
  
tokenized_dataset = dataset.map(tokenize_function, batched=True)
```

Explanation:

Converts text into tokens for model training, with labels matching input_ids so that it can learn to generate the full sequence.

◆ **Cell 10: Define Training Arguments**

```
training_args = TrainingArguments(  
    output_dir="./tinylama-qlora-support-bot",  
    per_device_train_batch_size=2,  
    gradient_accumulation_steps=4,  
    learning_rate=2e-4,  
    logging_dir="./logs",  
    num_train_epochs=3,  
    logging_steps=10,  
    save_total_limit=2,  
    save_strategy="epoch",  
    bf16=True,  
    optim="paged_adamw_8bit"  
)
```

Explanation:

Sets training config:

- 3 epochs
 - 2 batch size
 - Gradient accumulation for effective larger batch size
 - 8-bit optimizer (paged_adamw_8bit)
 - BF16 support if available (saves memory)
-

◆ **Cell 11: Start Training**

```
trainer = Trainer(  
    model=model,  
    args=training_args,  
    train_dataset=tokenized_dataset,  
    tokenizer=tokenizer  
)
```

```
trainer.train()
```

Explanation:

Creates a Trainer object and begins model fine-tuning on the tokenized dataset.

Cell 12: Save the Fine-Tuned Model

```
model.save_pretrained("tinylama-qlora-support-bot")  
tokenizer.save_pretrained("tinylama-qlora-support-bot")
```

Explanation:

Stores your trained model and tokenizer locally for future use or deployment.

Cell 13: Test the Model Locally

```
from transformers import pipeline
```

```
pipe = pipeline("text-generation", model=model, tokenizer=tokenizer)
```

```
instruction = "how can i request refund?"
```

```
prompt = f"### Instruction:\n{instruction}\n\n### Response:\n"
```

```
output = pipe(prompt, max_new_tokens=100)
```

```
print(output[0]['generated_text'])
```

Explanation:

Creates a test pipeline to generate responses for a sample instruction using the fine-tuned model.

◆ **Cell 14: Deploy as Gradio Chat UI**

```
import gradio as gr
```

```
def generate_response(instruction):
```

```
    prompt = f"### Instruction:\n{instruction}\n\n### Response:\n"
```

```
    output = pipe(prompt, max_new_tokens=100, do_sample=True,
temperature=0.7)
```

```
    return output[0]['generated_text'].replace(prompt, "").strip()
```

```
gr.Interface(
```

```
    fn=generate_response,
```

```
    inputs=gr.Textbox(lines=3, placeholder="Ask your customer support question here..."),
```

```
    outputs=gr.Textbox(lines=6),
```

```
    title="🔧 Customer Support Chatbot (TinyLlama + QLoRA)",
```

```
    description="Ask any support question. Model trained on MakTek/Customer_support_faqs_dataset using TinyLlama 1.1B."
```

```
).launch()
```